

A common application of graphs is to model some sort of transportation network (airports, roads, etc.).


If we are currently at a location s and wish to travel to a location t , then we may want to find a path which starts at s and ends at t . There may be many such paths, and in the application, some paths may be much more expensive to follow than others.

We can assign a weight to each edge $\{u, v\}$ of the graph which represents the cost of moving from location u to location v . Now consider some path p from s to t . The weight of the path $w(p)$ is the sum of the edge weights along the path.

$$p_1: s \rightarrow 2 \rightarrow 3 \rightarrow 4 \rightarrow t$$

$$w(p_1): 2 + 5 + 2 + 6 = 15$$

Example:




$$p_2: s \rightarrow 1 \rightarrow 3 \rightarrow 5 \rightarrow t$$

$$w(p_2) = 3 + 1 + 4 + 1 = 9$$


In this setting, we would be interested in computing a **shortest path** from s to t . A shortest path is a path of minimum weight from s to t . Let $\delta(u, v)$ denote the weight of a shortest path between any two vertices u and v in the graph ($\delta(u, v) = \infty$ if there are no paths from u to v).

In some applications we may want to have negative weights on an edge. Note that if there is a negative-weight cycle, then some shortest paths may not exist.



The following theorem is crucial in the computation of shortest paths:


Thm A subpath of a shortest path is a shortest path.



If this were not a shortest path from u to v , then we could replace this subpath with the shortest path from u to v and get a shorter path from s to t , a contradiction.

The following theorem is known as the *triangle inequality* and is also important in the computation of shortest paths:

Thm: For all $u, v, x \in V$ we have

$$\delta(u, v) \leq \delta(u, x) + \delta(x, v).$$



We will now consider the *single-source shortest path* problem in which we are given a graph with a designated source vertex s , and we wish to compute the shortest path weights $\delta(s, v)$ for each $v \in V$.

We will assume all edge weights are nonnegative so that we will not have any negative weight cycles.

Dijkstra's Algorithm. Idea: Greedy. For each vertex v , we maintain an upper bound $d[v]$ on $\delta(s, v)$.

1. Maintain a set S of vertices whose shortest path weights from s are known, that is $d[v] = \delta(s, v)$.
2. At each step, add the vertex $v \in V \setminus S$ whose $d[v]$ is minimal.
3. Update $d[u]$ for any vertex u adjacent to v .

Example of Dijkstra's Algorithm:



Correctness of Dijkstra's:

Theorem: (i) For all $v \in S : d[v] = \delta(s, v)$. (ii) For all $v \notin S : d[v]$ is the weight of a shortest path from s to v that uses only vertices in S (besides v itself).




The implication of this theorem is that Dijkstra's algorithm terminates with $d[v] = \delta(s, v)$ for each $v \in V$ (because each $S = V$ at the end of the algorithm).

The proof is by induction. It is clearly true in the base case ($d[s] = 0$ and $d[v] = \infty$ for all $v \neq s$). Assume (i) and (ii) are true before an iteration, and we will show it remains true after another iteration.

Let u be the vertex added to S in this iteration. So $d[u] \leq d[v]$ for all $v \in V \setminus S$.

(i) We need to show $d[u] = d(S, u)$.

Suppose the contrary. Then there is a path p from S to u with $w(p) < d[u]$. So there must be some vertex $y \in V \setminus S$ on the path p .

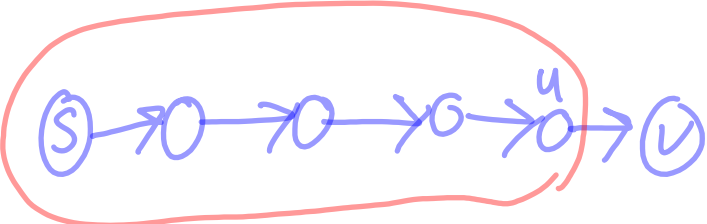


So we should have taken y instead of u , so such a path does not exist. Thus $d[u] = d(S, u)$.

(ii) Let $v \notin S$. Let p be a shortest path from S to v that only uses vertices in S (besides v itself).

2 cases:

1) p does not contain u . Then (ii) is true by inductive hypothesis.

2) p contains u : 

p consists of vertices in $S \setminus \{u\}$, then to u , then to v .

$\Rightarrow w(p) = d[u] + w(u, v)$ which is
we set $d[v]$ to after this iteration.

So (ii) is true.

The running time of Dijkstra's algorithm is $O(m \log n)$ when maintaining $V \setminus S$ as a priority queue.

Now consider the *unweighted case* in which we want to find a path with the smallest number of edges.

Certainly Dijkstra's Algorithm can still work (we can just set the weight of each edge to be 1). Can we do better?

Idea: do a modified BFS search starting from s . Recall BFS traverses the graph in “layers” where each layer will be the same distance from s .

So far we have only been computing the *length* of a shortest path. What if we want to know what the shortest path actually is?

We can build a shortest path tree similarly to how we built MST, BFS, and DFS trees (we remember the “predecessor” for each vertex during the computation).

